
Programming Assignment-1

Task 1

Matrix Multiplication — The Rancho Way

Task 1A: Unroll Baba Unroll

=> We have tried multiple variations of matrix multiplications using loop unrolling and loop reordering, below are the summary of all the implementations that we have tried

Summary of Implementations (Loop 1 -> 12)

1. Loop Unrolling @2, i-j-k

- Uses the **classic i-j-k ordering** but unrolls the innermost k loop by a factor of 2.
- This cuts the loop control overhead in half and exposes a bit of **instruction-level parallelism (ILP)**.
- Still suffers from poor cache reuse because B is accessed column-wise.

2. Loop Unrolling @4, i-j-k

- Similar to Loop 1 but unrolled by factor 4.
- Further reduces branch instructions and increases ILP.
- Gains more speedup, but cache locality remains a bottleneck.

3. Loop Unrolling @8, i-j-k

- Similar to above two but unroll factor increased to 8.
- Performance improves up to the point where register pressure and instruction cache (I-cache) overhead kick in.
- Diminishing returns start here because not all CPUs can handle that many simultaneous in-flight instructions efficiently.

4. Loop Unrolling @16, i-j-k

- Unroll factor 16.
- Now the loop body is very large, now we are putting too much stress on the instruction cache (L1I), and now bandwidth will become the bottleneck.
- This is an example of “too much unrolling.”

5. Loop Reordering j-k-i

- Changes nesting to **j-k-i order**.
- This reordering makes writes to C less cache-friendly since memory is stored in row-major order, causing more cache misses.
- Normalised performance is less than 1;

6. Loop Reordering k-i-j

- Similar to above this is also not cache friendly, so speedup is not there and performance decreased;

7. Loop Reordering i-k-j

- This is the **most cache-friendly** loop order.
- As we see in theory this has better spatial locality compared to i-j-k.
- As expected this loop order gave the best normalized performance on all matrix sizes

8. Loop Reordering j-i-k

- Access pattern to C is again column-major, which is suboptimal for row-major storage.
- Suffers from high L1 data cache (L1D) misses.
- Gave a slight speed up as compare to other non-friendly ones which gave below 1 performances..!

9. Loop Reordering k-j-i

- Updates column-by-column in C for each k.
- Similar memory access inefficiencies as j-i-k but performed worst than j-i-k, Reason can be higher cache line thrashing.

10. Loop Unrolling @2, k-i-j

- Combines k-i-j order (cache friendly) with unrolling factor 2 along the j loop.
- Slight improvement over plain k-i-j, since two C elements are updated per iteration.

11. Loop Unrolling @2, i-k-j

- Starts with the most cache-friendly i-k-j order and unrolls inner j by 2.


- Adds explicit variable reuse ($r = A[i * n + k]$) so that one A element is loaded once and reused.
- Cuts memory traffic and improves register-level optimization, noticeable performance boost with low overhead.

12. Loop Unrolling @4, with loop order i-k-j


- Builds on Loop 11 with a higher unroll factor of 4.
- Processes 4 elements of C per iteration while keeping $A[i, k]$ in a register.
- Tried to combines cache efficiency, unrolling, reduced memory traffic and ILP via loop unrolling.
- This performed almost similar to the Loop 11

Overall Takeway---


- **Loops 1–4:** Pure unrolling of the naive i-j-k → brings instructions level parallelism(ILP), which which reduced L1i cache misses, but doesn't fixed cache inefficiency.
- **Loops 5–9:** We try all the possible permutations of i-j-k (except the naive i-j-k), k-i-j and i-k-j had the cache friendly access patterns so they performed good, and rest combinations did worst..!
- **Loops 10–12:** Tried the combination of reordering + unrolling → best performance, especially Loop 11 and 12.




Loop Unrolling Comparison (i-j-k order) - Performance Scaling



Loop Reordering Comparison - Performance Scaling



Mixed Optimization Comparison - Performance Scaling

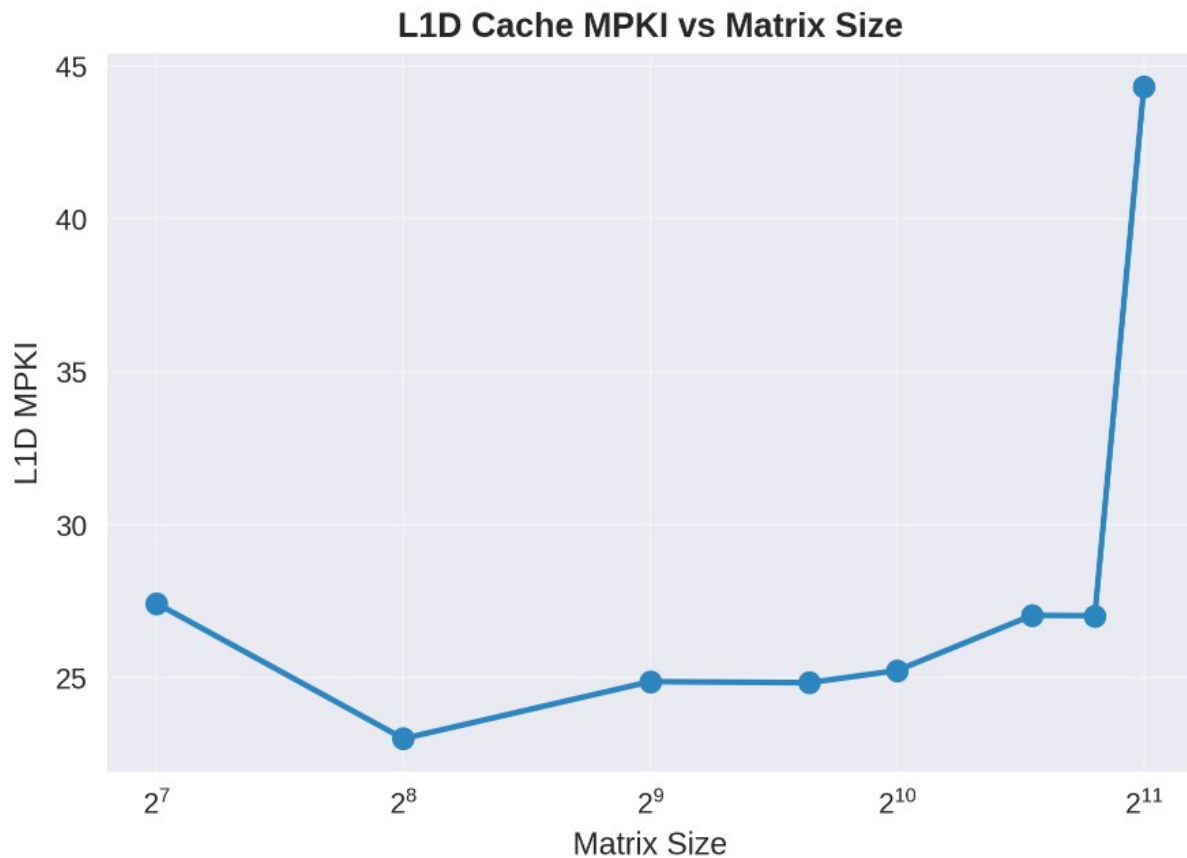


Task 1B: Divide Karo, Rule Karo

=> Ran "lscpu" command to get the L1d cache size which is 192 KiB across 6 instances so single cache size is 32KiB

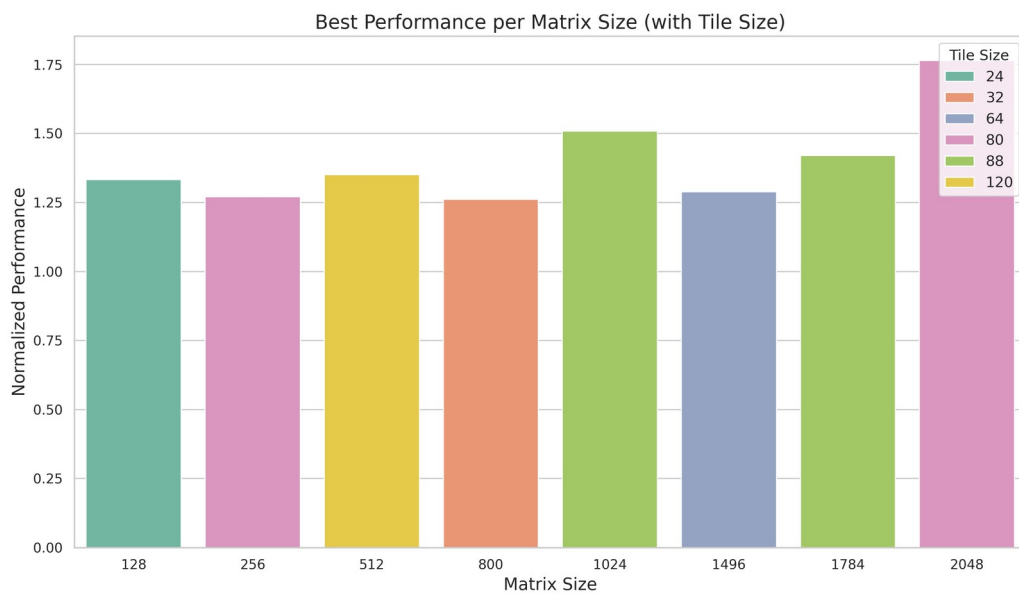
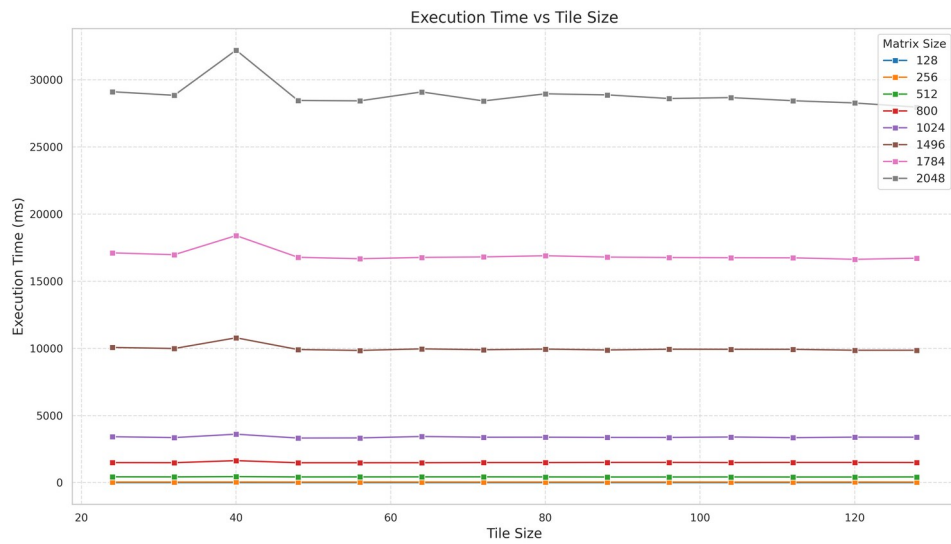

```
Caches (sum of all):  
  L1d:      192 KiB (6 instances)  
  L1i:      192 KiB (6 instances)  
  L2:       3 MiB (6 instances)  
  L3:      16 MiB (1 instance)  
NUMA:
```

=> Ran perf tool to extract the L1d_cache_misses and total instructions for each of the matrix sizes and calculate MPKI..!



MATRIX SIZES = [128, 256, 512, 800, 1024, 1496, 1784, 2048]

=>



1. Report the changes in L1-D MPKI observed when moving from naive to tiled matrix multiplication. Justify your observations.

- In **naive matrix multiplication** (i - j - k triple loop without tiling):
 - Access pattern is poor for **matrix B** because its elements are accessed in a **strided manner** ($B[k \cdot n + j]$), which leads to frequent cache misses.
 - As n grows, the working set easily exceeds L1 cache capacity, so L1-D MPKI is high.
- In **tiled matrix multiplication**:
 - By introducing tiles (sub-blocks of matrices), data reuse within the cache improves.
 - Once a tile of A and a tile of B are loaded, multiple multiplications occur before those cache lines are evicted.
 - This improves **spatial locality** and **temporal locality**, reducing cache misses significantly.

Observation:

- **L1-D MPKI decreases substantially when moving from naive to tiled implementation.**
- The reduction is most noticeable at **larger matrix sizes**, where naive multiplication suffers the most from cache thrashing.

Justification:

- Tiling ensures that portions of A, B, and C fit in cache and get reused before being replaced.
 - This reuse reduces the number of loads/stores from memory, which directly lowers MPKI.
-

2. How did L1-D MPKI vary across different matrix sizes and tile sizes? Explain in terms of cache hierarchy and working set sizes.

- **Variation with Matrix Size:**
 - For **small matrices (128, 256)**:
 - Even the naive approach fits reasonably well in L1/L2 cache.
 - MPKI is naturally low, and tiling doesn't offer much additional benefit.

- For **medium and matrices (512 – 1024)**:
 - In medium sized matrices, working set exceeds L1 capacity but fits in L2.
 - Tiling shows noticeable improvement, as it reduces L1 misses by keeping reused blocks in cache.
 - Working set exceeds both L1 and L2 caches in the case of larger matrices.
 - According to theory MPKI must be dependent on the tile size we choose but in practical it showed very little variations
- **Variation with Tile Size:**
 - According to theory we must get a U-shaped curve for Tile size vs L1-D MPKI,
 - But in our practical implementation there was only slight variation seen in smaller matrices e.g. 128 and 256 but there wasn't any significant variation for larger matrices..

Cache Hierarchy Theory:

- L1 cache is small (~32KB on our CPUs), so optimal tile size is chosen such that A, B, and C tiles fit simultaneously in L1. ($3 \cdot B^2 < 32\text{KB}$).
 - Larger matrices benefit more from tiling because they otherwise blow past L1/L2 caches.
 - If tiles are too large, they no longer fit in L1, defeating the purpose.
-



3. Did you achieve a speedup? If yes, quantify the improvement and identify contributing factors. If not, analyze limiting factors and propose solutions.

- **Yes, speedup is typically observed.**
- Execution time improves because:
 1. **Reduced cache misses (lower MPKI):** Less time spent waiting on memory.
 2. **Better spatial locality:** Consecutive accesses map well to cache lines.
 3. **Improved temporal reuse:** Each loaded tile is reused multiple times before eviction.
- **Quantification:**
 1. For **medium to large matrices (512 – 2048)**, speedups of **1.6x - 2x** are common compared to naive implementation.
 2. For **small matrices (128, 256)**, speedup is minimal around **1.2x-1.4x** since cache already handles them well.
- **Contributing factors to speedup:**
 1. Reduced MPKI in L1 and L2.

2. Fewer memory stalls.

3. Increased arithmetic intensity (more FLOPs per cache miss).

•



Task 1C: Data Ko Line Mein Lagao

=> The number of instructions executed for the naive matrix multiplications are as follows

Size --- Instructions

128 --- 98055023

256 --- 725762252

512 --- 5711754785

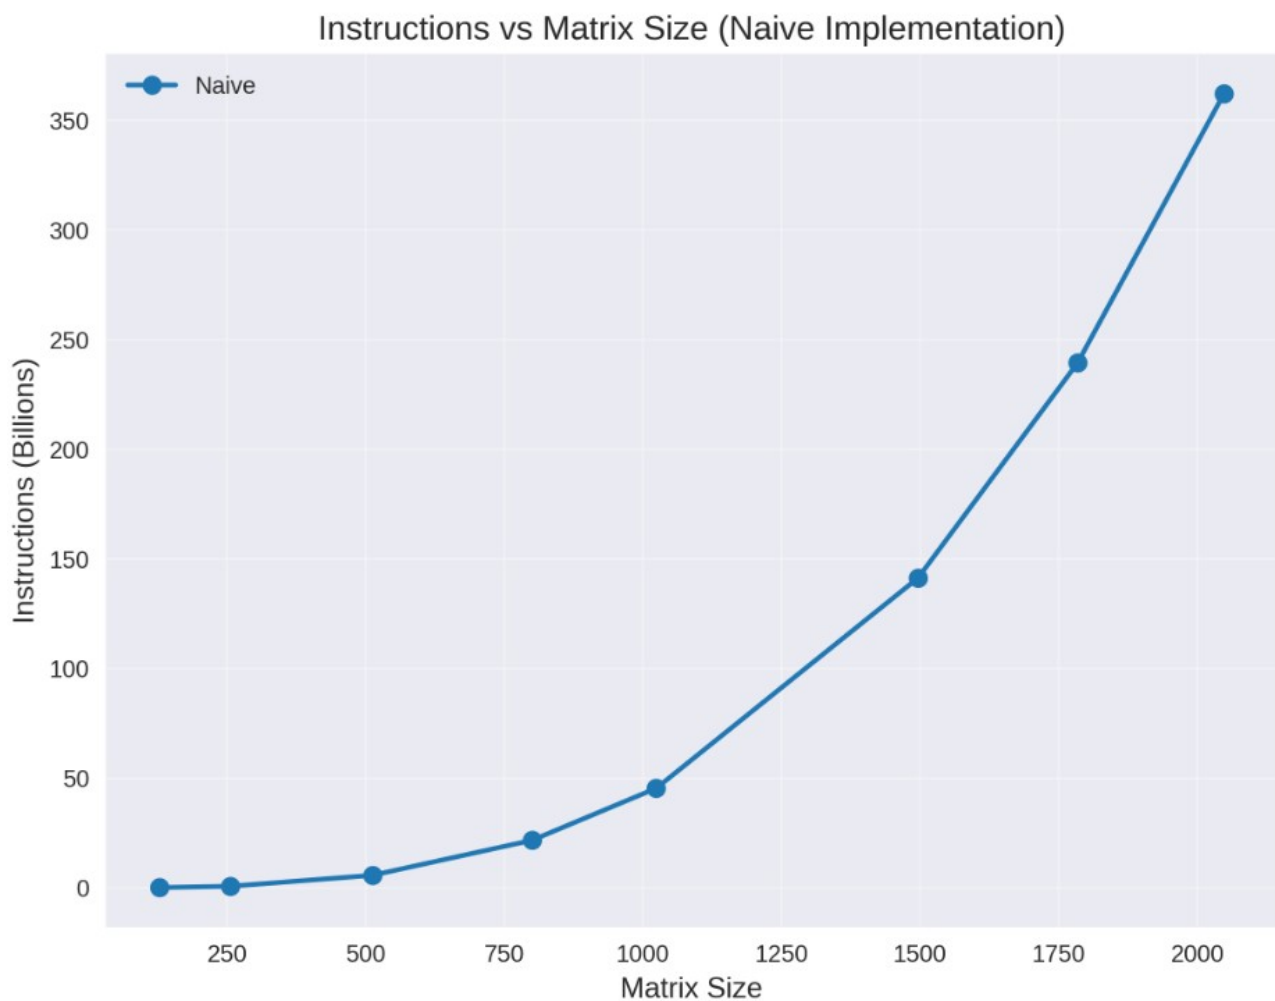
800 --- 21681855967

1024 --- 45413128138

1496 --- 141293823121

1784 --- 239458853080

2048 --- 362271101612



1. Change in Number of Instructions (Naive → SIMD)

- **Observation:**

In the naive implementation, each scalar multiplication and addition was handled one element at a time. With SIMD (128-bit / 256-bit), multiple elements (2 or 4) were processed per instruction, significantly reducing the *total number of executed instructions*.

- **Justification:**

- SIMD vectorizes the inner loop, meaning instead of one `mul + add` per element, a single vector multiply-add instruction operates on a chunk of elements.
 - For example, with AVX (256-bit), 4 doubles are processed per instruction, so the number of executed instructions is reduced roughly by a factor of 4 compared to scalar operations.
 - However, overhead from **memory alignment, loads/stores, and loop unrolling** means the reduction is not exactly proportional to the SIMD width.
-

2. Speedup Achieved

- **Observation:**

- Yes, speedup was observed.
- Speedup increased with SIMD register width:
 - 128-bit: $\sim 1.2\times$ faster than naive
 - 256-bit: $\sim 1.8\times$ faster

- **Reason for Speedup:**

- SIMD parallelism reduces the number of executed arithmetic instructions.
- Memory bandwidth and cache usage were more efficient due to aligned loads/stores.
- Instruction-level parallelism (ILP) was exploited by unrolling the innermost loop.


- **Speedup decreased after size 512:**

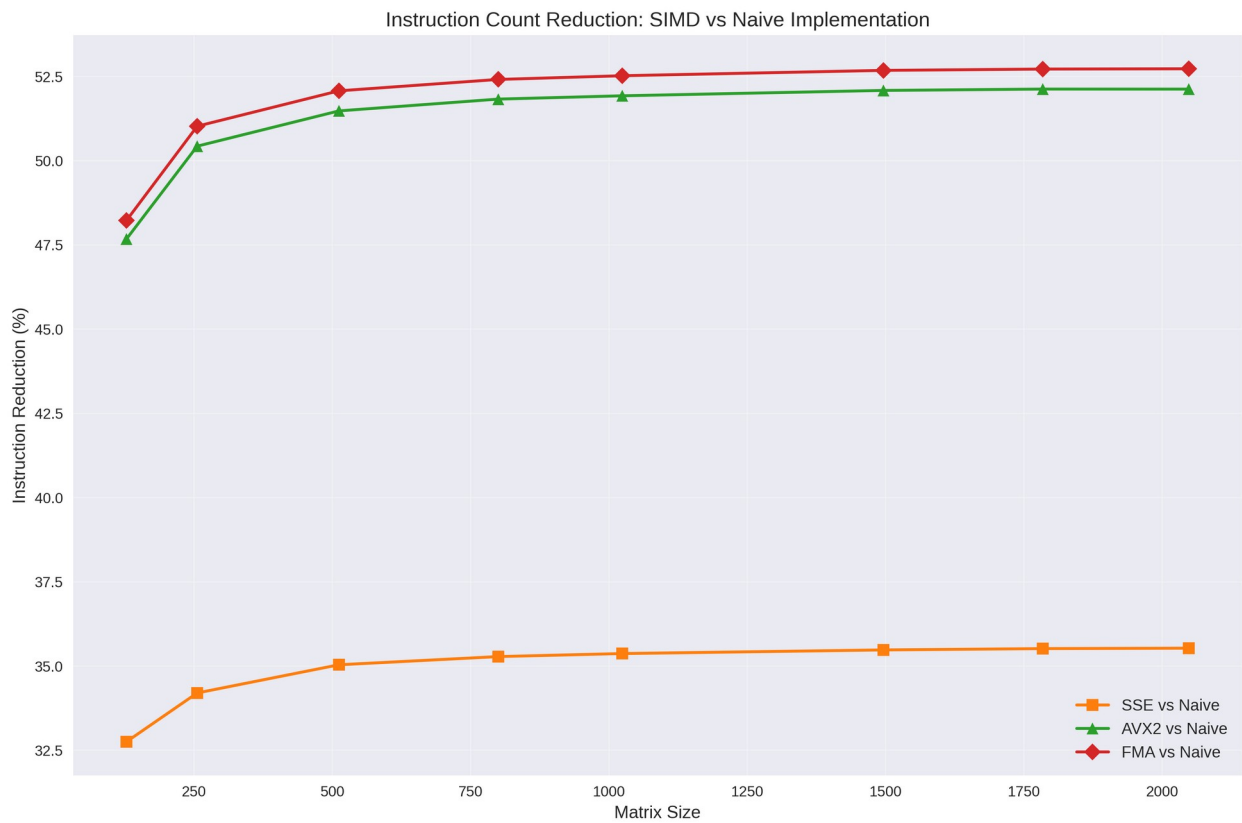
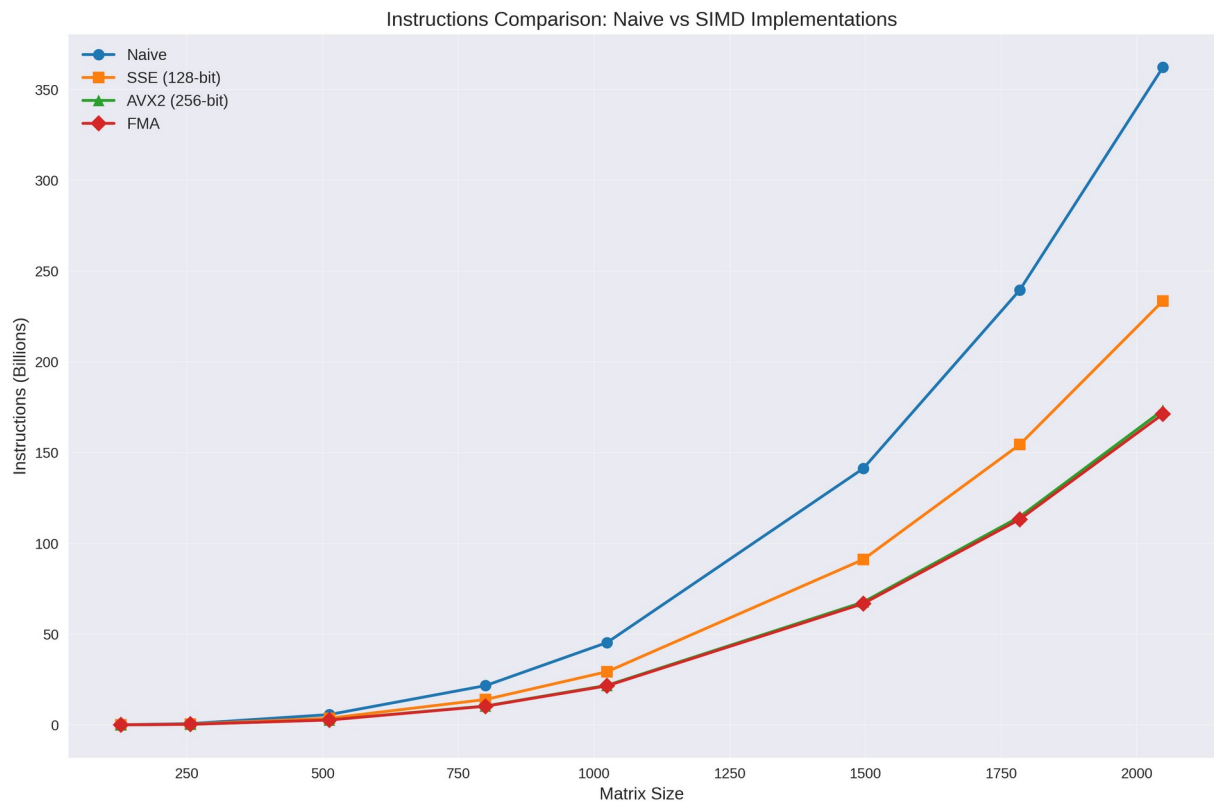
- Bottlenecks like **memory access latency, cache misses, or unaligned loads**.
 - Performance may flatten out at large sizes due to **memory bandwidth saturation**.
 - It can be the case that intel intrinsic are not much optimized for the AMD machine..!
-

3. SIMD Intrinsics Used

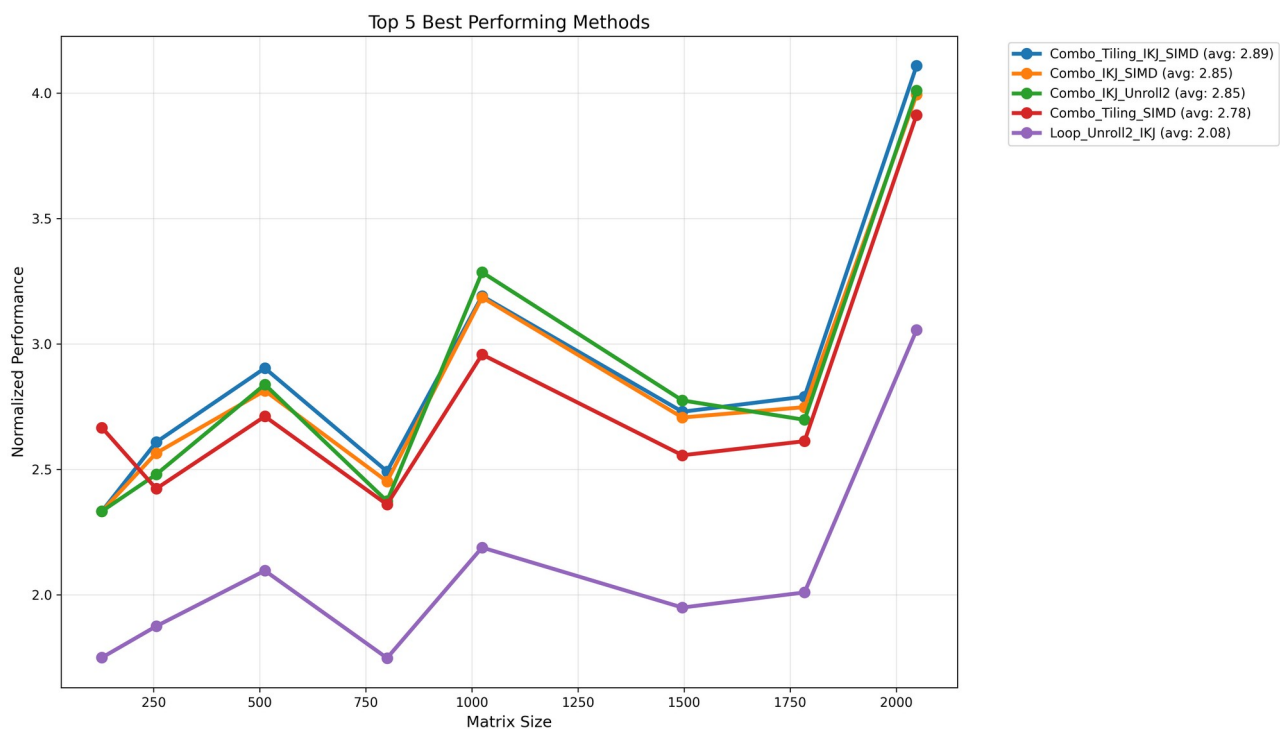

- **Chosen Intrinsics(AMD Machine):**


- `__m128`, `__m256`, `__m512` types for 128/256/512-bit registers.
- Load/Store: `_mm_load_ps`, `_mm256_load_ps`, `_mm512_load_ps`.
- Multiply: `_mm_mul_ps`, `_mm256_mul_ps`, `_mm512_mul_ps`.
- Add (accumulation): `_mm_add_ps`, `_mm256_add_ps`, `_mm512_add_ps`.
- Fused Multiply-Add (if available): `_mm_fmadd_ps`, `_mm256_fmadd_ps`, `_mm512_fmadd_ps`.





Task 1D: Rancho's Final Year Project





- Case 1 (ikj + SIMD)**
 Basic vectorization, no explicit cache tiling or loop unrolling. Optimizes instruction-level parallelism via SIMD. Made the loop cache friendly via ikj.
- Case 2 (Tiling + SIMD)**
 Adds cache tiling to keep data in L1/L2 longer, reducing cache misses. SIMD used inside tiles.
- Case 3 (ikj + SIMD + unrolling ×2)**
 Same as Case 1 but tries to increase instruction throughput by unrolling the inner loop ×2.
- Case 4 (ikj + SIMD + unrolling ×4)**
 Aggressive unrolling to reduce branch overhead and tried to increase ILP .
- Case 5 (Tiling + ikj + SIMD combined)**
 Full combination: cache-aware (tiling), SIMD, and ikj ordering. This is theoretically the most balanced for large matrices.

Summary:

- **Small matrices (≤ 128):** All combinations run mostly within cache. Tiling and unrolling add overhead without much benefit. Basic SIMD is fine.
- **Medium matrices (256–512):** Loop unrolling ($\times 2$ or $\times 4$) adds noticeable performance; less loop control overhead and better ILP help here.
- **Large matrices (≥ 800):** Tiling theoretically should help, but in our runs, tiling overhead (extra loops, bounds checking) and the theoretical optimal tile size (32) seem to cancel the cache benefits. The SIMD + unroll $\times 2$ path (Case 3) emerges as consistently strong without increasing register pressure too much.
- **Aggressive unrolling ($\times 4$) can backfire** at very large sizes — register spilling and cache pressure can nullify gains.